



Applied Machine Learning in Engineering

Lecture 03 winter term 2023/24

Prof. Merten Stender

Cyber-Physical Systems in Mechanical Engineering, Technische Universität Berlin

www.tu.berlin/cpsme

merten.stender@tu-berlin.de

Recap: Lecture 02



		Type	Description	Example	Operations
categorical / qualitative	{	Nominal	value corresponds to a set of mutually exclusive values, classes or categories	eye color, city name, brand name, class name, booleans	= and \neq
		Ordinal	values from a set of distinct values, can be put into an order	house numbers, study semester, grades	= and \neq ($<$, \leq , $>$, \geq)
numeric / quantitative	{	Interval	values from a continuous set of equally spaced values, unit of measurement exists, no true zero	temperature $^{\circ}\text{C}$, time on 12-hour clock, IQ test results	= and \neq ($<$, \leq , $>$, \geq) ($+$, $-$)
		Ratio	values from a continuous set of equally spaced values, unit of measurement exists, true zero indicates absence	age, velocity, height	= and \neq ($<$, \leq , $>$, \geq) ($+$, $-$, $*$, $/$)

Recap: Lecture 02



- Making qualitative (categorical) data readable to a computer
- **One-Hot Encoding** $y \in \{v_1, \dots, v_k\} \mapsto y \in \mathbb{R}^k, \in \{0, 1\}, k: \text{number of distinct values / classes}$

- Traffic example:
 - 4 classes: pedestrian, crosswalk, stop sign, car

- pedestrian $\rightarrow [1 \ 0 \ 0 \ 0]^T$
- crosswalk $\rightarrow [0 \ 1 \ 0 \ 0]^T$
- stop sign $\rightarrow [0 \ 0 \ 1 \ 0]^T$
- car $\rightarrow [0 \ 0 \ 0 \ 1]^T$

- Model prediction $y = [0.95 \ 0 \ 0 \ 0.05]^T \rightarrow \text{to 95\% a pedestrian, to 5\% a car}$

Do not use for ordinal data, as order gets lost!

(almost) no ML algorithm does create an implicit relationship or order between neighboring values in an output array such as a OHE vector



Integer Encoding

- Encoding ordinal data requires **keeping an order**
- Simplistic encoding: assign an integer to each category, start with 0 for the first category
- **Example:** satisfaction rating for this class
 - 5 classes with natural rank order (*“extremely dislike”, “dislike”, “neutral”, “like”, “extremely like”*)
 - Integer encoding:

extremely dislike	→	0
dislike	→	1
neutral	→	2
...		
- **Caution!** Integer encoding keeps the order, but pretends a measure of (equal) distances!
 - Strictly speaking, a model prediction $\tilde{y} = 1.8$ is not meaningful, and rounding may be wrong
 - Decoding strategy is highly case-specific!

Recap: Exercise 02



- Dictionaries and mappings between key-value pairs
- One-hot encoding at set of categorical values

```
# obtain number of classes and classes themselves
self.classes = np.unique(values)
self.num_classes = len(self.classes)

# mapping between category (key) and index (value) via dictionary
self._class_map = dict(zip(self.classes, np.arange(stop=self.num_classes)))

# one-hot encode the categorical data
encoded_vals = []
for val in values:
    _enc_value = np.zeros(self.num_classes) # empty vector of zeros
    _enc_value[self._class_map[val]] = 1 # turn 'hot' (put in a one)
    encoded_vals.append(_enc_value) # stack to existing values
```

Recap: Exercise 02



```
def decode(self, enc_vals):
    """ Invert one-hot encoding.
    expects a binary N x K binary matrix. N samples, K categories
    returns list or one-dimensional array of categories
    """
    # inverse mapping between index and category
    self._inv_class_map = {v: k for k, v in self._class_map.items()}

    # de-code one-hot encoded values
    values = []
    for enc_val in enc_vals:

        idx = np.argwhere(enc_val == 1)[0][0]
        values.append(self._inv_class_map[idx])

    return np.hstack(values) # return a one-dim. np.ndarray of categories
```

```
if __name__ == "__main__":

    """
    Testing the implementation
    """

    OHE = OneHotEncoder()

    values = np.array(['Berlin', 'Frankfurt', 'Munich', 'Berlin'])
    print(f'values to encode: \t{values}')

    # fit the encoder
    OHE.fit(values)

    ohe_values = OHE.encode(values)
    print(f'one-hot encoded representation: \n {ohe_values}')

    dec_values = OHE.decode(ohe_values)
    print(f'de-coded values: \t{dec_values} \n\n')
```

Agenda



Cyber-Physical Systems
in Mechanical Engineering TU Berlin

- Unsupervised learning
- Cluster validity metrics
- K-means clustering
- Python: scikit-learn library

Learning outcomes



Cyber-Physical Systems
in Mechanical Engineering TU Berlin

Learn to ...

- Identify unsupervised learning tasks
- Quantify properties of clusters
- Implement a basic clustering method

Know about ...

- Time complexity of K-means
- Advantages and disadvantages of K-means



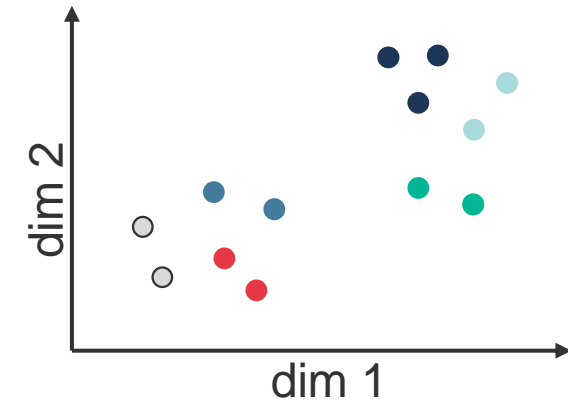
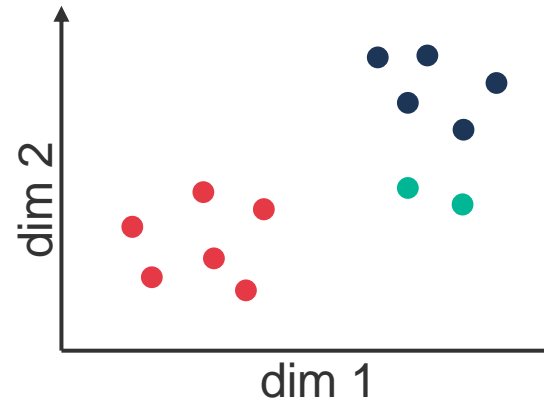
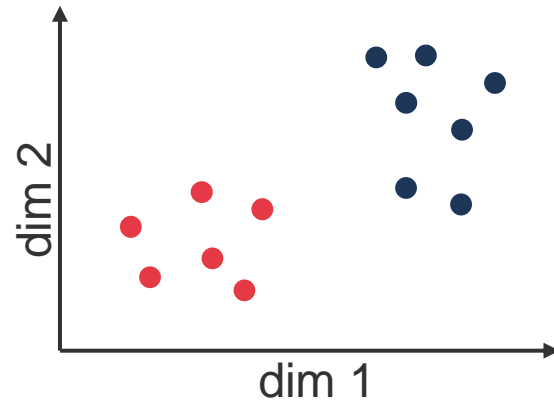
Clustering

Clustering Data



Clustering denotes the process of dividing a set of data points into distinct groups (clusters), thereby **maximizing intra-group similarity** and **minimizing inter-group similarity**.

- What to consider a good clustering?



→ Measures of cluster validity

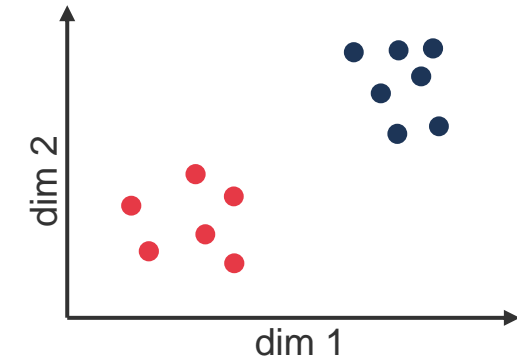
Measures of Cluster Validity



The goodness of a clustering is in the eye of the beholder

- **Well-separated** clusters denote a situation in which any point in a cluster is closer to every other point in the cluster than to any point not belonging to the cluster

- **Objectives for cluster validity measures:**
 1. Avoid finding patterns in noise
 2. Create robust, repeatable and consistent clusterings
 3. Find a meaningful number of clusters
 4. Maximize similarity inside clusters and maximize difference between clusters



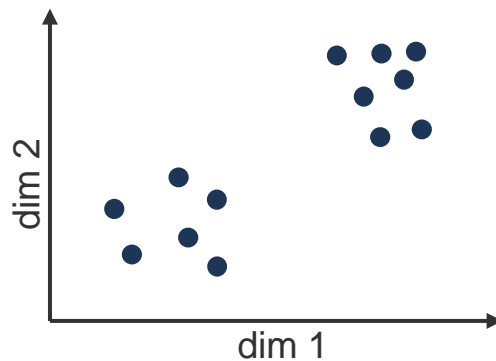
Measures of Cluster Validity



Internal measures

No a-priori information exists about clusters or class labels

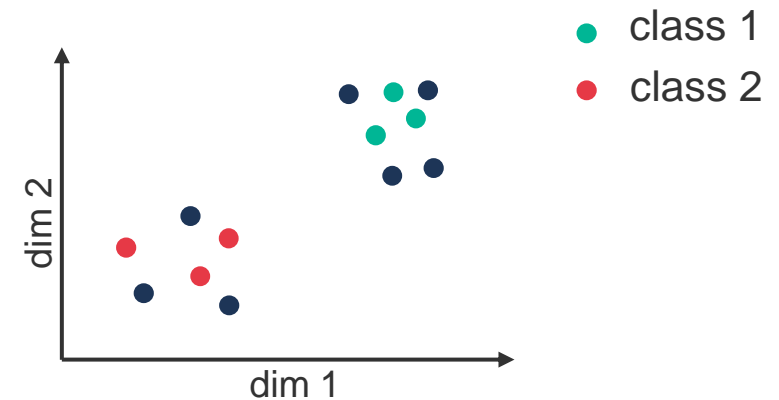
- Momentum (SSE/compactness/cohesion)
- Cluster separation
- Silhouette Coefficient
- Dunn-index, correlation, ...



External measures

Some external knowledge about clusters exist, such as class labels for some instances.

- Entropy
- Adjusted Mutual Information (AMI)
- Adjusted Rand Index (ARI), ...

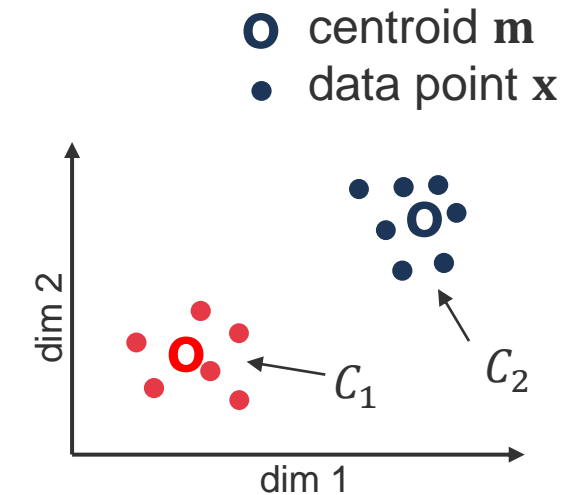


Momentum (Compactness / Cohesion)



- Internal cluster validity metric based on **within-cluster sum-of-squares**

$$\text{SSE} = \sum_{k=1}^K \sum_{\mathbf{x}_i \in C_k} \|\mathbf{m}_k - \mathbf{x}_i\|^2 \quad \text{centroid } \mathbf{m}_k \in \mathbb{R}^n$$



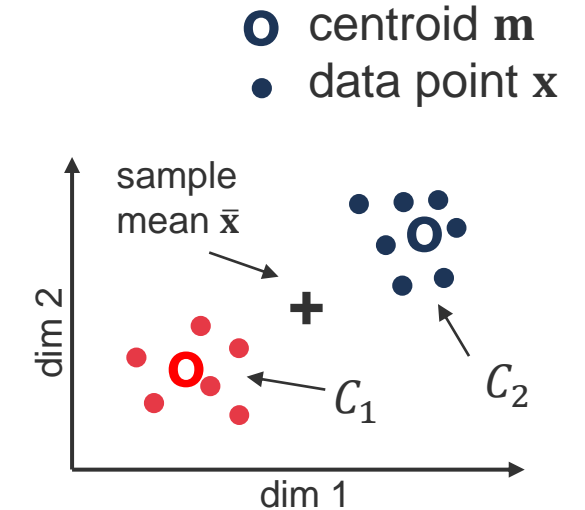
- Measure of the **variability** and **density** of the observations within each cluster
- Question:** what would the clustering look like if we optimized for minimal SSE?
 - w/o constraints on the number of clusters: $\text{SSE} = 0$ for $K = N \rightarrow$ undesired behavior
 - Simplest / trivial approach: set user-defined number of clusters $K \rightarrow$ **K-means algorithm**

Cluster Separation



- Internal cluster validity metric based on **between cluster sum of squares**

$$\text{BSS} = \sum_{k=1}^K \|\mathbf{m}_k - \bar{\mathbf{x}}\|^2 \quad \text{centroid } \mathbf{m}_k \in \mathbb{R}^n$$



- Measure for how far centroids are spread out w.r.t. sample mean $\bar{\mathbf{x}}$
- BSS does not account for cluster size, cluster density, or well-separated clusters
- Larger BSS is better



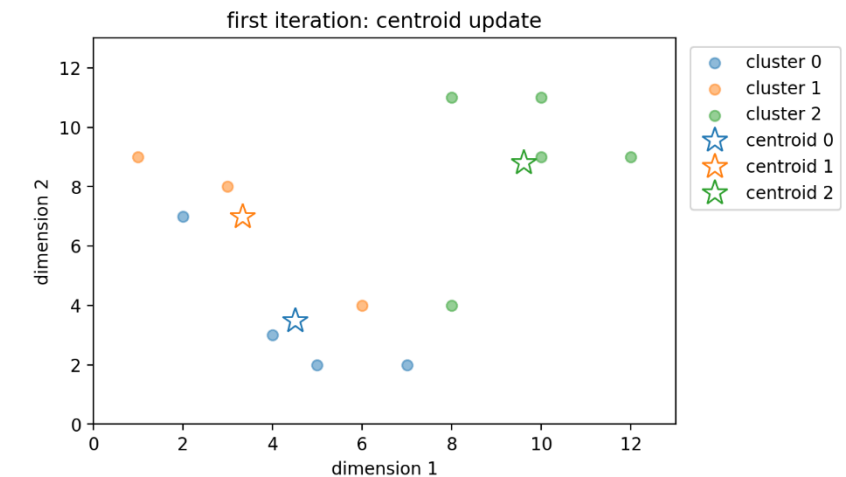
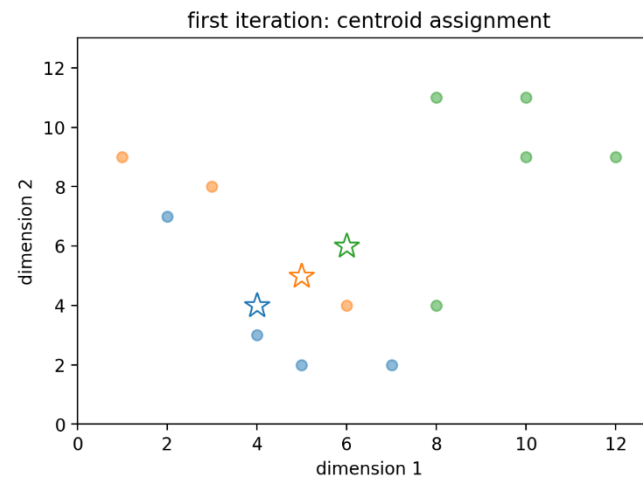
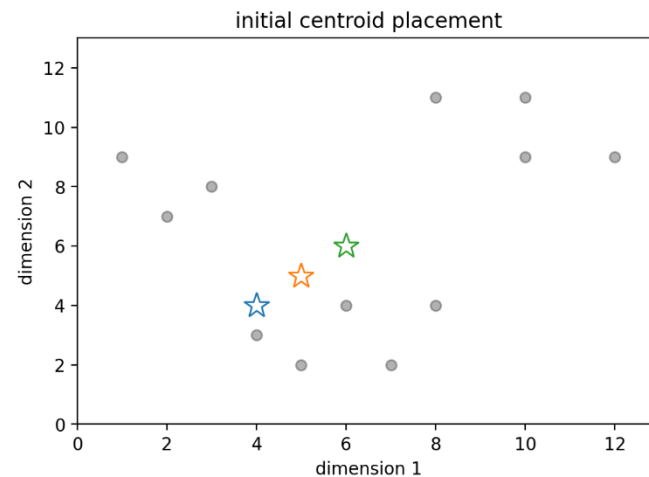
K-means clustering

K-means: Basic Algorithm



- Simplest and very efficient clustering algorithm
- Prototype-based, finding K (user-defined) clusters by optimal centroid placement

1. Placement of K random centroids
2. Loop until converged:
 1. Assign data points \mathbf{x}_i to closest centroid \mathbf{m}_k to build cluster \mathcal{C}_k
 2. Update centroid position by averaging across $\mathbf{x}_i \in \mathcal{C}_k$

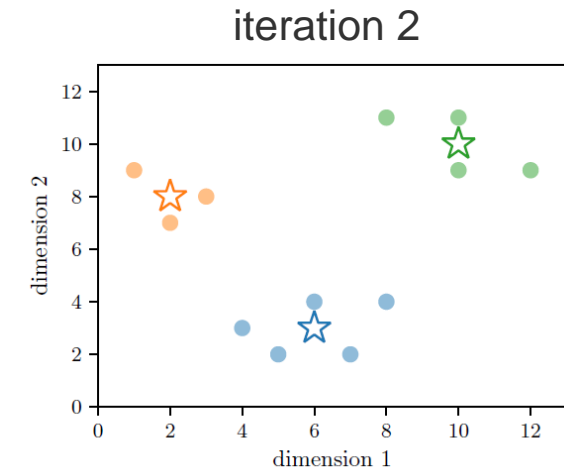
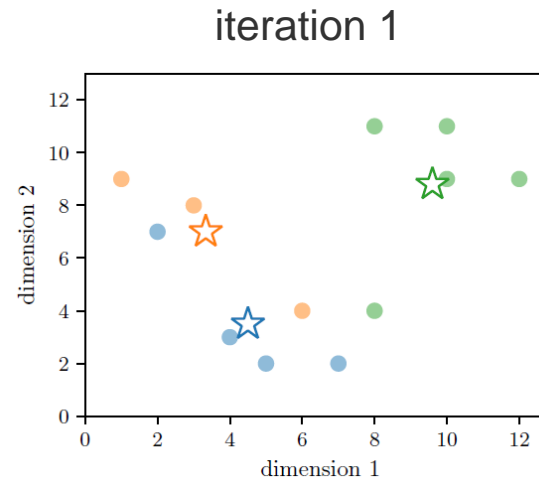
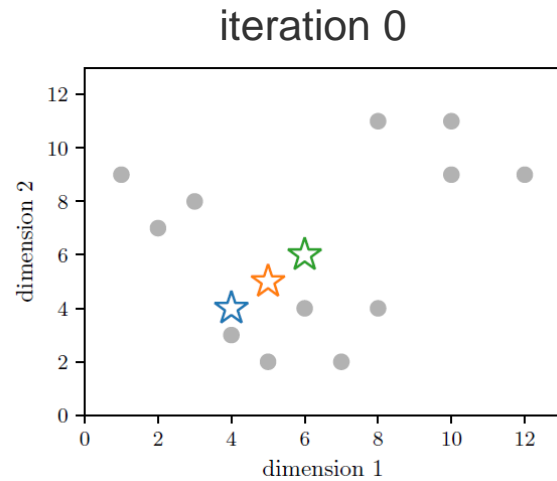


K-means: Convergence Criterion

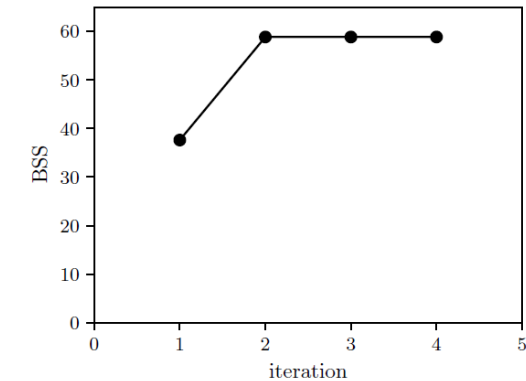
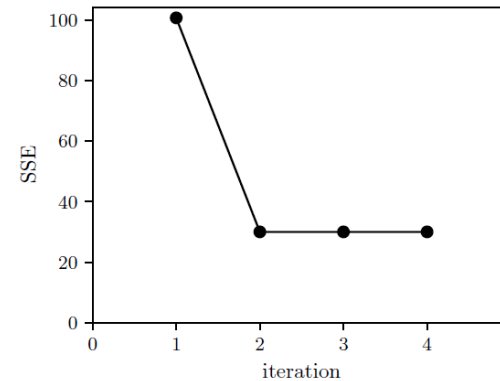


- **When to stop updating centroid assignments and centroid updates?**
- Consider some consecutive iterations of K-means
 - How many points changed clusters?
 - How much did centroid positions change? → Frobenius norm of consecutive $\mathbf{m} = [\mathbf{m}_1, \dots, \mathbf{m}_K]$
- Typical convergence criterion: <1% points change clusters during last 2 iterations

K-means: Example



- Tracking cluster validity metrics



K-means: Complexity



- Simple and efficient algorithm

Algorithm 1 *K*-means algorithm (basic form)

```
1: Select  $K$  initial centroids  $\mathbf{m}_k$ 
2: repeat
3:   for all  $i = 1, \dots, N$  do                                ▶ First phase: cluster assignment
4:     assign point  $\mathbf{x}_i$  to closest centroid  $\mathbf{m}$  and update assignment vector  $r_{ik}$ 
5:   end for
6:   for all  $k = 1, \dots, K$  do                                ▶ Second phase: centroid update
7:     update centroid positions  $\mathbf{m}_k$  by averaging across all  $\mathbf{x}$  in cluster  $C_k$ 
8:   end for
9: until clusterings converged
```

- **Complexity:** $O(K \cdot N \cdot n \cdot j)$ (scales **linearly** with number of data points)
(j : number of iterations until convergence)

K -means (basic): Pitfalls and Caveats



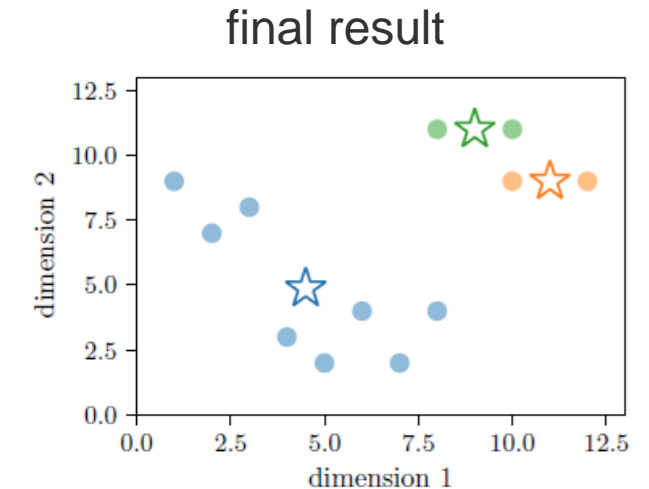
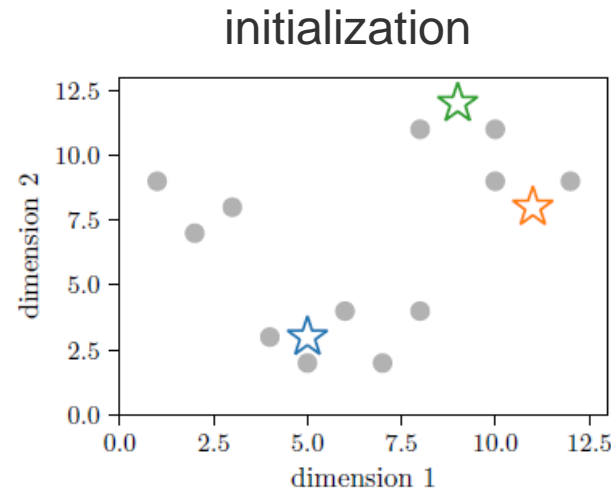
- Weak convergence (trapped in local minimum)
 - Strong dependence on initial centroids
- Empty clusters
 - Algorithm stops when a cluster is empty
- Non-deterministic results
 - For stochastic centroid initialization
- User-defined selection of K
 - Generally unknown, requires parameter studies
- Sensitivity to noise and outliers
 - Requires a-priori outlier detection or more advanced K -means algorithms

Dependence on Initial Centroids



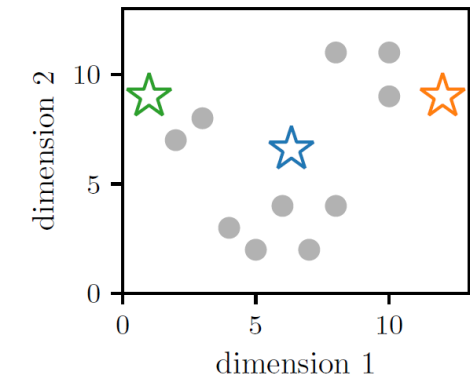
- Weak convergence for poor centroid initialization

- Centroids placed randomly in data range
- Centroids picked randomly from data points



- Solution strategies

- Repetitive clustering**, each time selecting different centroids, selecting minimal
- Iterative centroid placement**: 1st centroid into data sample center, 2nd into data point farthest away, 3rd centroid into data point farthest away from 1st and 2nd centroid, ...
- User-defined centroid placement** (leveraging a-priori knowledge)



Handling Empty Clusters



- Clusters C can end up empty after the cluster assignment step
- No point will ever be assigned to that cluster again
- Basic algorithm would stop once an empty cluster is met

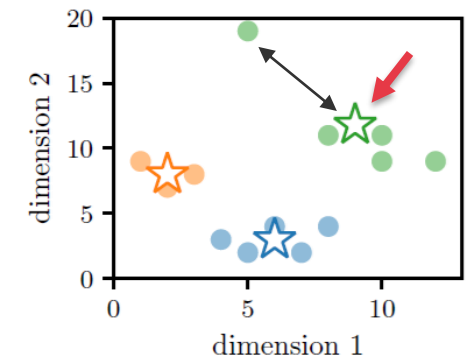
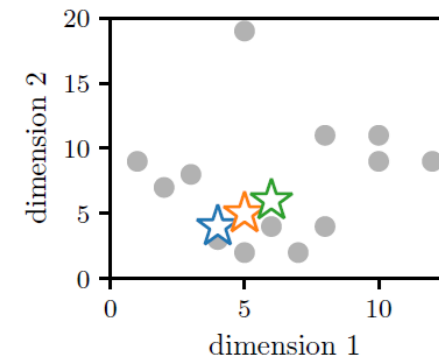
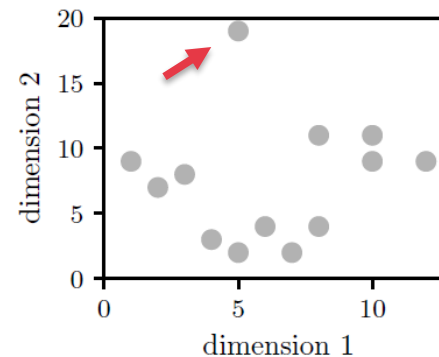
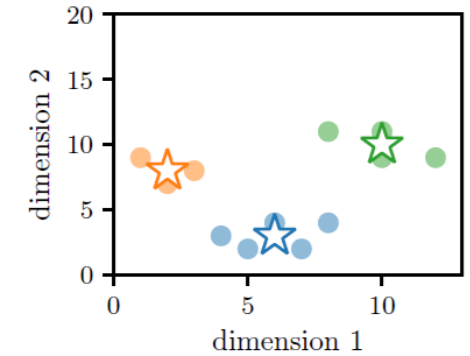
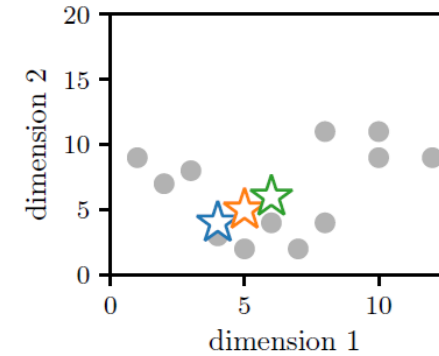
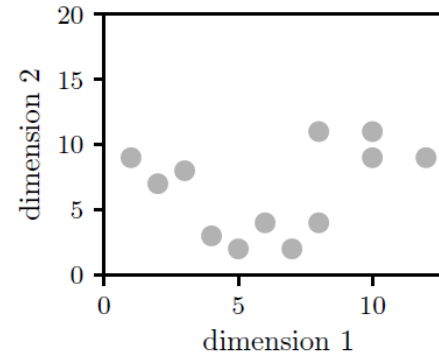
- Solution strategy: centroid re-placement
 1. Place centroid at position of data point the farthest away from any centroid
→ Maximal reduction of total SSE, prone to selecting outliers

 2. Place centroid into the cluster that is the less compact (largest SSE value)
→ Splits large clusters, potentially creating artificial sub-clusters

Sensitivity to Noise and Outliers



- Noise and/or outliers can heavily distort the final centroid positions
- Outliers: largest contribution to cluster validity metrics (SSE)

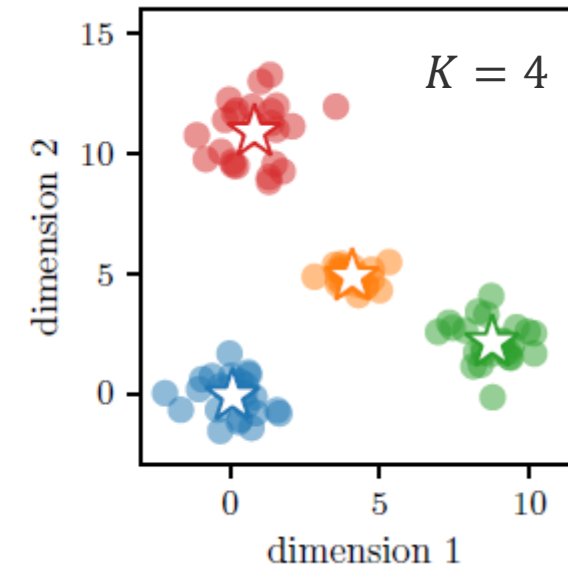
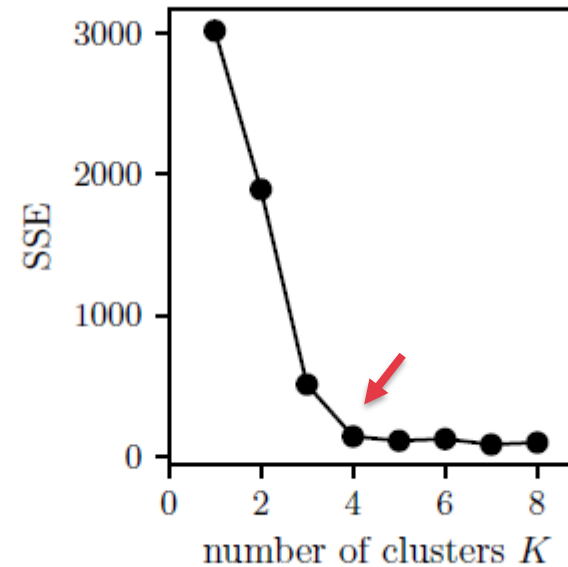
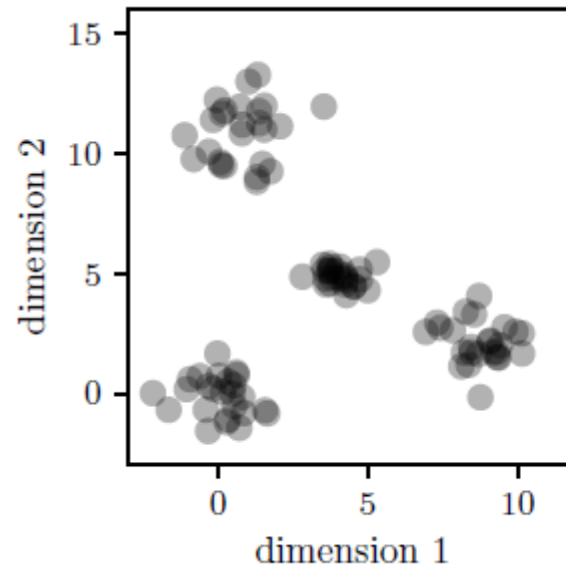


- Strategies
 - Outlier removal before clustering
 - Advanced K -means methods: remove extraordinarily strongly contributing data points
 - Post-processing by data point removal and re-running K -means from there

Selecting an Optimal K



- Hyperparameter study for K while tracking cluster validity metrics
- Selection of final K : elbow method



Variants of K -means



- Today: **Lloyd's algorithm** (simplest and basic form)

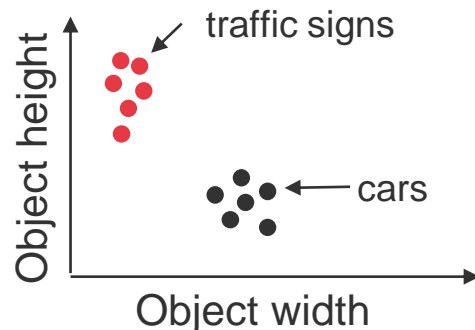
Other approaches:

- **MacQueen's algorithm**: updates centroid positions with each new data point assignment (incremental updating)
 - Better convergence behavior
 - Introduces an order-dependency (no determinism can be installed), non-repetitive results
- **Hartigan-Wong's algorithm**: initial centroids placed in vicinity of data center; centroid assignment based on SSE or other cluster validity metric instead of (Euclidean) distance metric
 - Designed to build more compact clusters
 - Initialization prone to generating artificial subclusters
- **Many more**: bisecting K -means (better convergence), ...

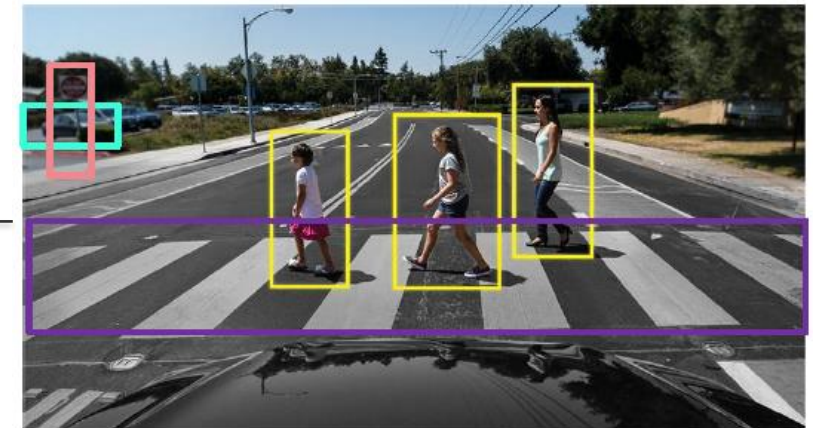
Unsupervised Learning: Applications



- Find similar customer behavior: shopping carts at the supermarket
 - Potential clusters: students, young family, retiree
- Data (image) compression and color quantization ([link](#))
- Engineering applications:
 - Finding similar customer behavior, such as driving conditions
 - Finding patterns in high-dimensional measurement data
 - Data pre-processing for ML modeling

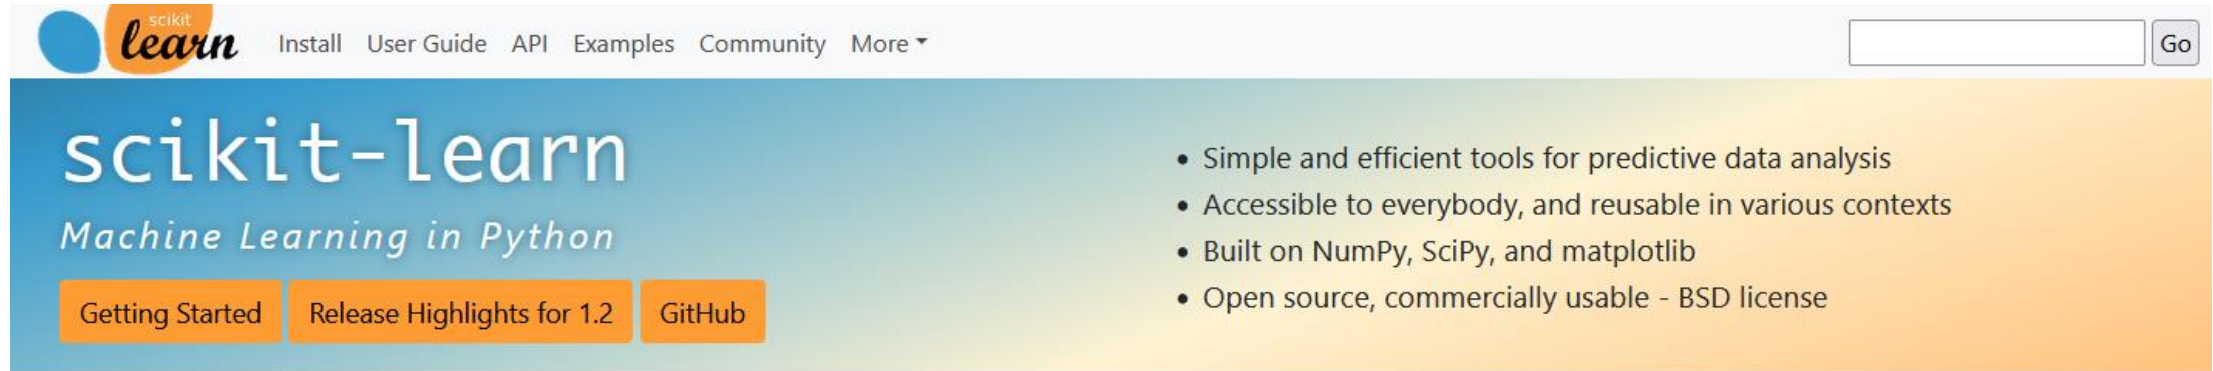


Anchor definition for object detection methods





Python: scikit-learn




- Most-used Python-based package for classical (statistical) machine learning
- Well-documented with underlying mathematics and literature
- Common structure across all ML methods: `.fit()`, `.predict()` methods
- K-means clustering: <https://scikit-learn.org/stable/modules/generated/sklearn.cluster.KMeans.html>

K-means clustering in scikit-learn



- `sklearn.cluster.KMeans` ([link](#))

 [Install](#) [User Guide](#) [API](#) [Examples](#) [Community](#) [More](#)

[Prev](#) [Up](#) [Next](#)

scikit-learn 1.2.2
[Other versions](#)

Please [cite us](#) if you use the software.

[sklearn.cluster.KMeans](#)
[KMeans](#)
[Examples using sklearn.cluster.KMeans](#)

sklearn.cluster.KMeans

```
class sklearn.cluster.KMeans(n_clusters=8, *, init='k-means++', n_init='warn', max_iter=300, tol=0.0001, verbose=0, random_state=None, copy_x=True, algorithm='lloyd')
```

[\[source\]](#)

K-Means clustering.

Read more in the [User Guide](#).

Parameters:	
n_clusters : int, default=8	The number of clusters to form as well as the number of centroids to generate.
init : {'k-means++', 'random'}, callable or array-like of shape (n_clusters, n_features), default='k-means++'	Method for initialization: 'k-means++': selects initial cluster centroids using sampling based on an empirical probability distribution of the points' contribution to the overall inertia. This technique speeds up convergence. The algorithm implemented is "greedy k-means++". It differs from the vanilla k-means++ by making several trials at each sampling step and choosing the best centroid among them. 'random': choose <code>n_clusters</code> observations (rows) at random from data for the initial centroids. If an array is passed, it should be of shape (n_clusters, n_features) and gives the initial centers. If a callable is passed, it should take arguments X, n_clusters and a random state and return an initialization.



Exercise 03

November 8th, 2023

Exercise 03



- Implement K-means algorithm (template provided, some lines to add)
- Evaluate on a small sample data set
- Compare against scikit-learn implementation
- [Extra]: Implement the same functionality as object-oriented code



Questions?