



# Applied Machine Learning in Engineering

**Lecture 08, June 06, 2023**

Prof. Merten Stender

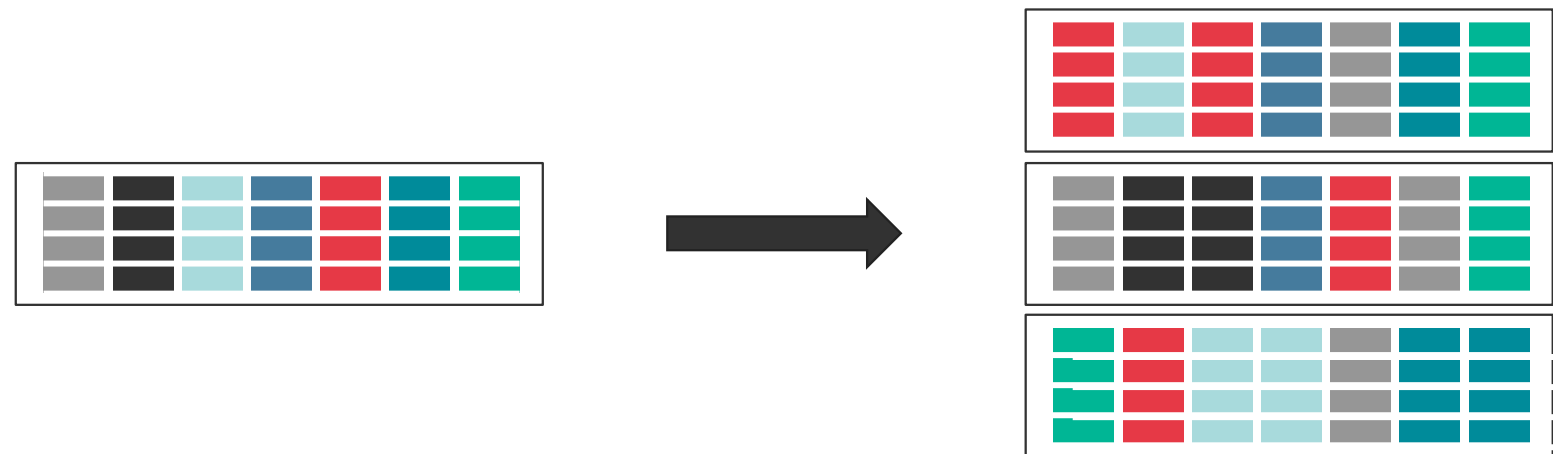
Cyber-Physical Systems in Mechanical Engineering, Technische Universität Berlin

[www.tu.berlin/cpsme](http://www.tu.berlin/cpsme)  
[merten.stender@tu-berlin.de](mailto:merten.stender@tu-berlin.de)

# Recap: Lecture 07



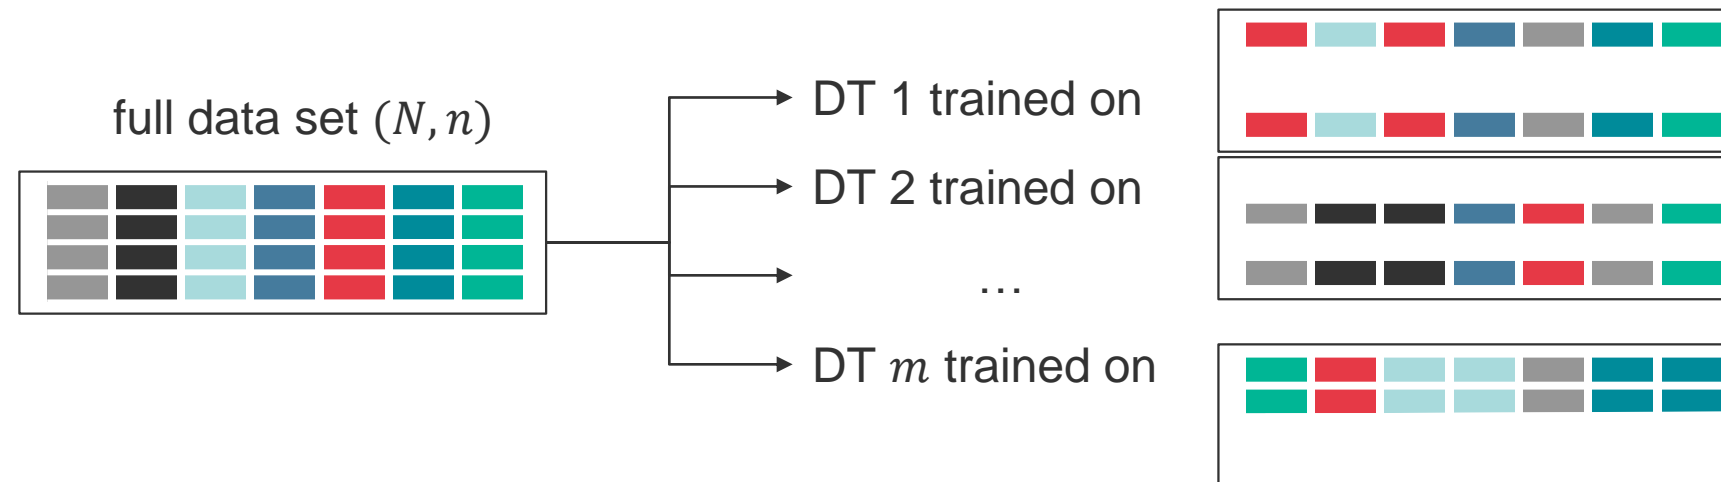
- **Ensemble methods:** combining weak learners to a stronger one
  - Bagging
  - Boosting
  - Random Forests
- **Bagging:** Training  $m$  models on  $m$  bootstrapped data subsets



# Recap: Random Forests



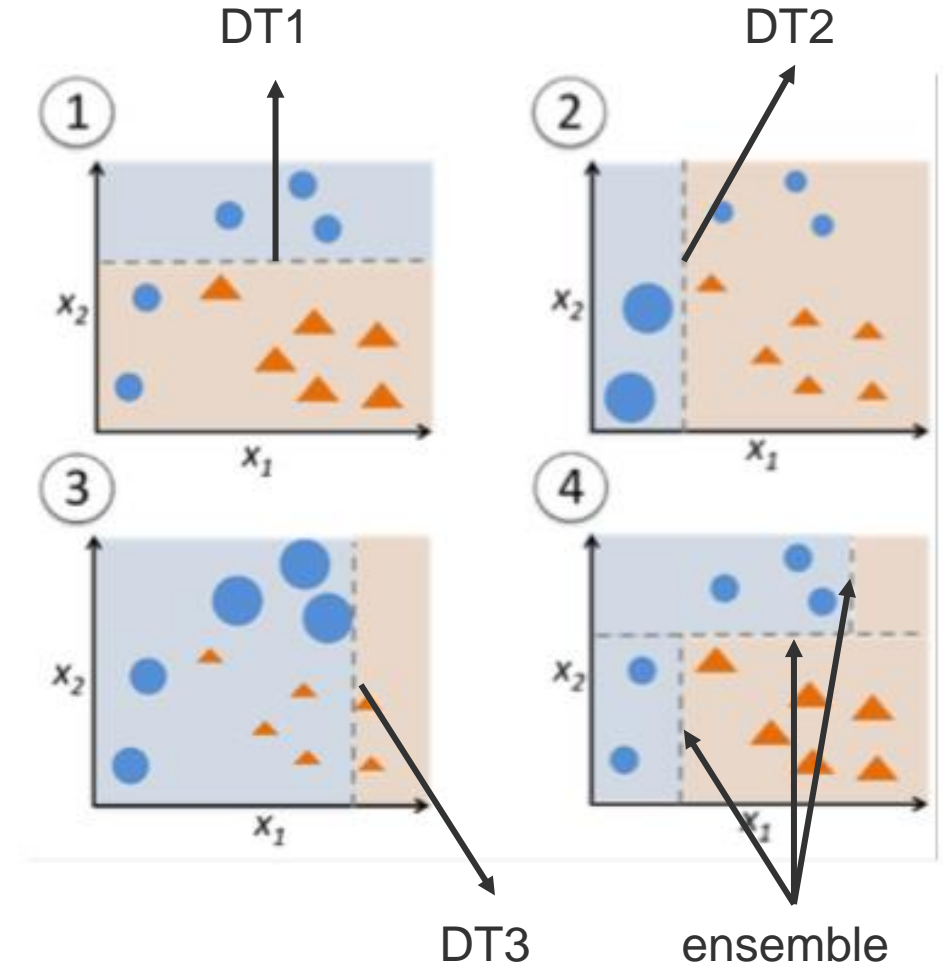
- Combination of bagging with random feature selection



# Recap: Boosting



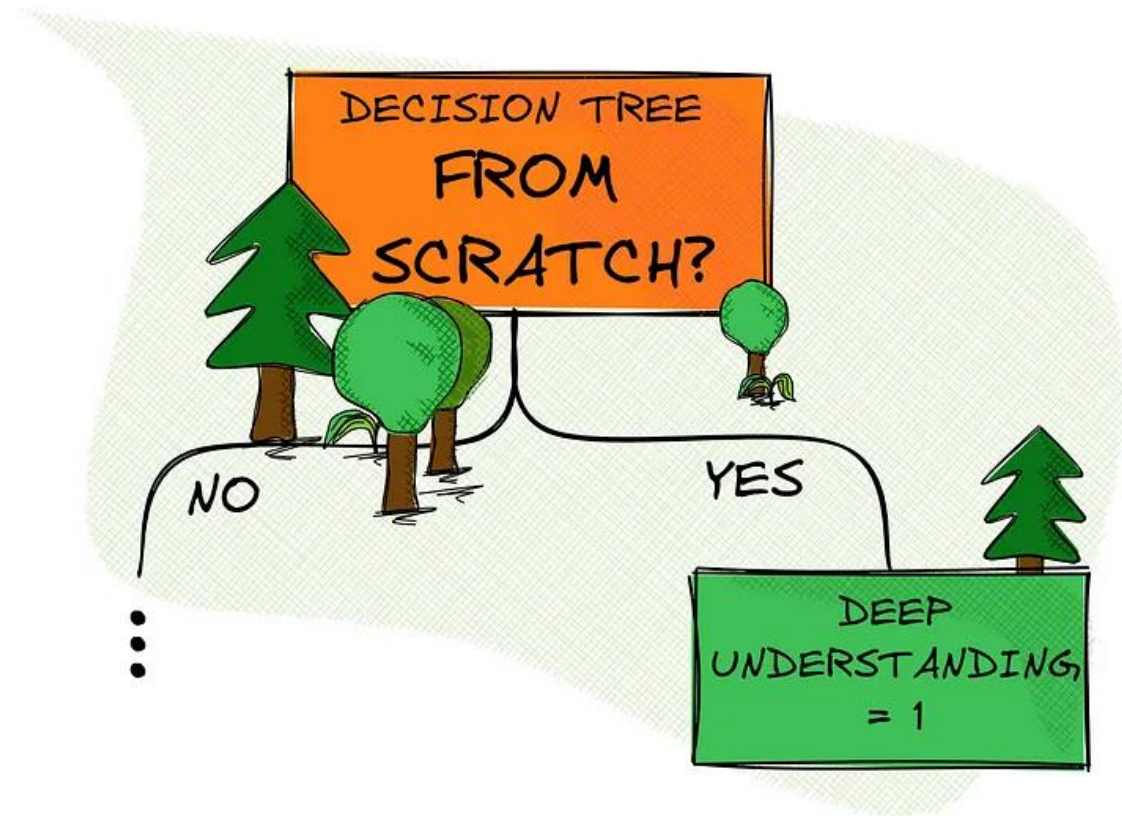
- Weak learners: decision tree (DT) with only 1 splitting rule (two leaf nodes)  
***decision tree stump model***
- Build first tree, find data set regime where it performs weak
  - Put larger weight on mis-classified data points
  - Build second tree using weighted data set
- Continue (without weight reduction)



# Recap: Exercise 07



- From scratch implementation of a decision tree for classification
- Recursive functions to grow the tree and traverse data through the tree



© Marvin Lanhenke, <https://towardsdatascience.com/implementing-a-decision-tree-from-scratch-f5358ff9c4bb>

# Recap: Exercise 07



```
# 0.0 base condition (stopping criteria) for recursive call.
# return leaf node if base condition is met
if self.is_completed(depth=curr_depth):
    return Node(prediction=self._majority_vote(y))

# 1. find best split for current data (X, y)
split_dim, split_val = self._best_split(X=X, y=y)

# 2. split the current data according to best split and
# return index for data sent to left and right child node
left_idx = self._create_split(X, split_dim, split_val)
right_idx = ~ left_idx

# 3. create child nodes and assign data using recursive call
child_left = self._grow_tree(X=X[left_idx:], y=y[left_idx], curr_depth=curr_depth+1)
child_right = self._grow_tree(X=X[right_idx:], y=y[right_idx], curr_depth=curr_depth+1)

# create a node with child nodes for the given split condition
node = Node(split_dim=split_dim, split_val=split_val, child_left=child_left, child_right=child_right)
```

# Today



Cyber-Physical Systems  
in Mechanical Engineering TU Berlin

- Recognize supervised classification tasks
- Understand  $k$ -nearest neighbors and logistic regression models
- Evaluate classification prediction models

# Agenda



Cyber-Physical Systems  
in Mechanical Engineering TU Berlin

## Machine Learning:

- K-nearest neighbors
- Logistic regression
- Classification metrics
- Data set splitting and k-fold cross validation

## Python:



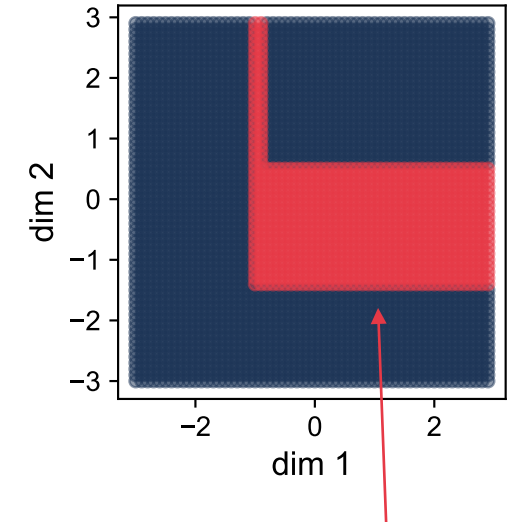
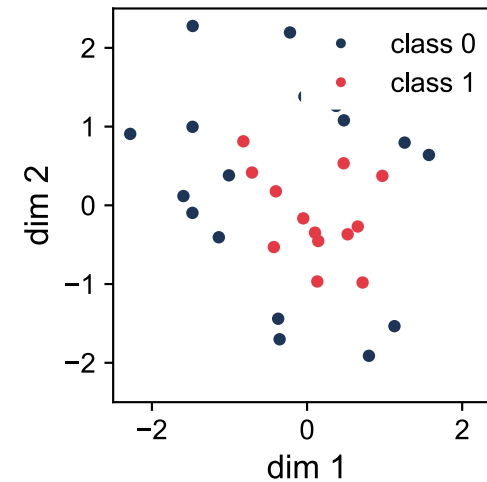


# k-nearest neighbors (kNN)

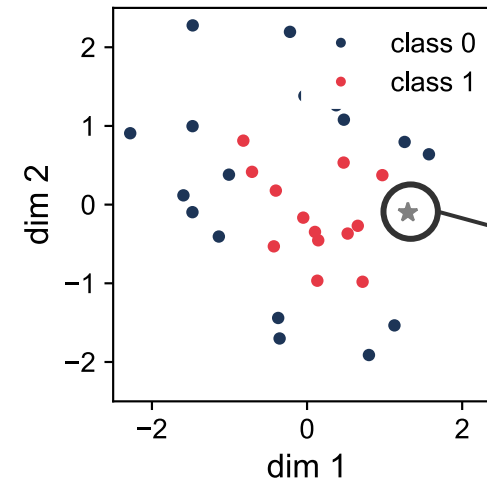
# Motivation



- Recap from exercise 07: binary classification in 2-dimensional space
- Decision Tree:
  - Classification of new, unseen data points based on binary feature space segmentation
  - Issue: overfitting
- Can you think of a simpler way to assign labels to new data points?



decision  
boundary of DT



which label?

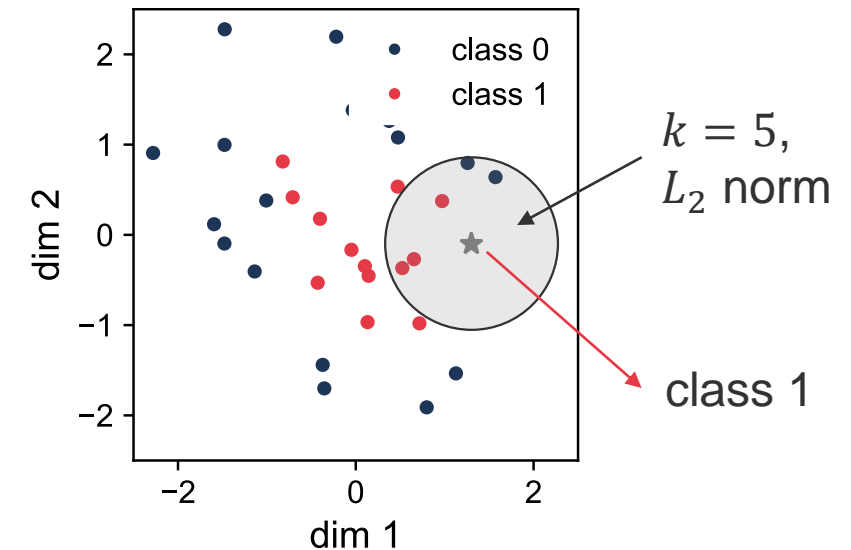
# $k$ -nearest neighbor algorithm (classification)



Cyber-Physical Systems  
in Mechanical Engineering TU Berlin

**Underlying assumption:** Two data points in close distance will be likely to belong to the same class

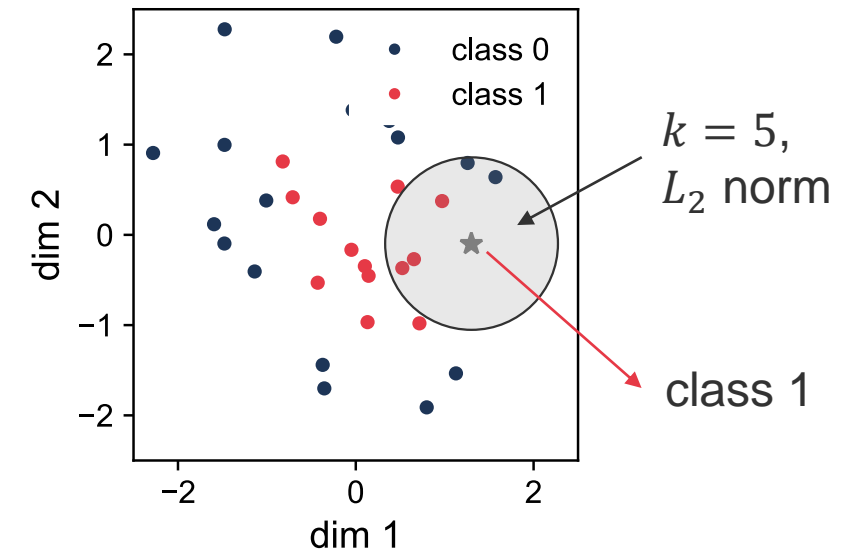
- **$k$ -nearest neighbor** (kNN): classification of new data point based on the majority vote across its  $k$  nearest neighbors w.r.t. norm  $\|\cdot\|$
- **Lazy learning:** Unparameterized model, no internal model parameters  $\theta$  are adjusted to training data samples. Training data is stored and queried for predictions on new data.
- **Hyperparameter:** number of neighbors  $k$ 
  - Empirical choice of  $k$  for balancing under- and overfitting



# $k$ -nearest neighbor algorithm



- **Regression task:** return mean / median target value of  $k$  nearest neighbors
- (Time) **complexity:**
  - $O(N)$  (trivial approach)
  - $O(\log N)$  efficient kd-tree neighborhood search
- **Limitations and caveats**
  - High-dimensional data sets: curse of dimensionality
  - Data on different scales: normalization required
- **Extensions:**
  - Probability prediction (classification): return probability instead of majority vote
  - Sample weights based on inverse of distance
  - Sample weights based on importance, e.g. based on trust in training data samples



# $k$ -nearest neighbor algorithm



## `sklearn.neighbors.KNeighborsClassifier`

```
class sklearn.neighbors.KNeighborsClassifier(n_neighbors=5, *, weights='uniform', algorithm='auto', leaf_size=30, p=2, metric='minkowski', metric_params=None, n_jobs=None) \[source\]
```

Classifier implementing the k-nearest neighbors vote.

Read more in the [User Guide](#).

### Parameters:

**`n_neighbors : int, default=5`**

Number of neighbors to use by default for `kneighbors` queries.

**`weights : {'uniform', 'distance'}, callable or None, default='uniform'`**

Weight function used in prediction. Possible values:

- 'uniform' : uniform weights. All points in each neighborhood are weighted equally.
- 'distance' : weight points by the inverse of their distance. In this case, closer neighbors of a query point will have a greater influence than neighbors which are further away.
- [callable] : a user-defined function which accepts an array of distances, and returns an array of the same shape containing the weights.

**`algorithm : {'auto', 'ball_tree', 'kd_tree', 'brute'}, default='auto'`**

Algorithm used to compute the nearest neighbors:

- 'ball\_tree' will use `BallTree`
- 'kd\_tree' will use `KDTree`
- 'brute' will use a brute-force search.
- 'auto' will attempt to decide the most appropriate algorithm based on the values passed to `fit` method.

Note: fitting on sparse input will override the setting of this parameter, using brute force.

**`leaf_size : int, default=30`**

Leaf size passed to `BallTree` or `KDTree`. This can affect the speed of the construction and query, as well as the memory required to store the tree. The optimal value depends on the nature of the problem.

**`p : int, default=2`**

Power parameter for the Minkowski metric. When  $p = 1$ , this is equivalent to using `manhattan_distance` (L1), and `euclidean_distance` (L2) for  $p = 2$ . For arbitrary  $p$ , `minkowski_distance` (L<sub>p</sub>) is used.

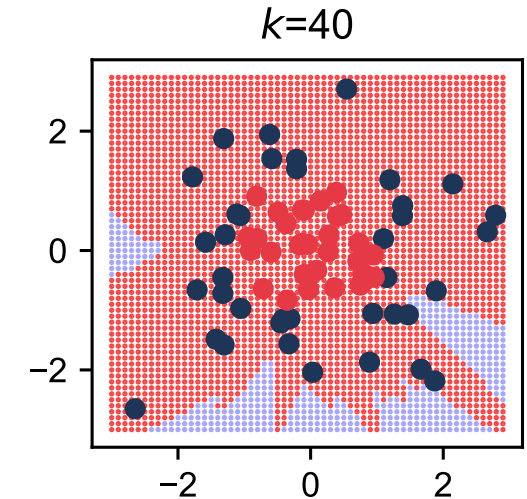
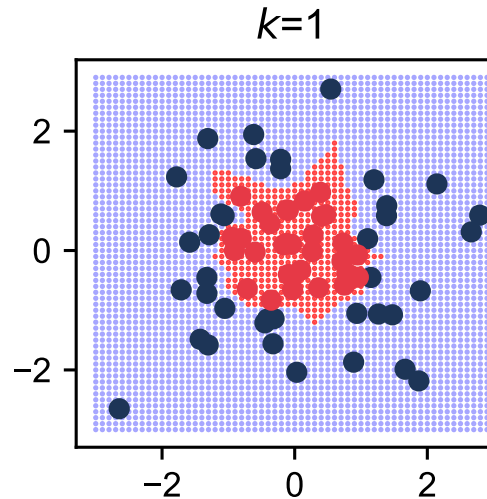
# kNN: choosing $k$



- Hyperparameter  $k$ : number of nearest neighbor to consider

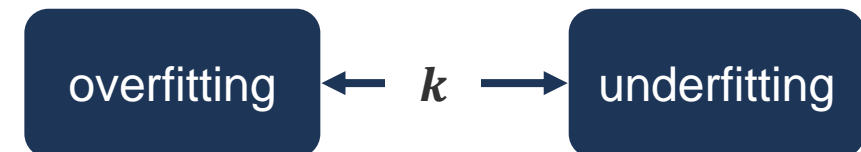
- (Too) small  $k$ :

- High accuracy on training data
- Low accuracy on validation data
- Complicated decision boundary
- Risk of overfitting



- (Too) large  $k$ :

- Low accuracy on training and validation data set
- Blurred decision boundaries
- Risk of underfitting

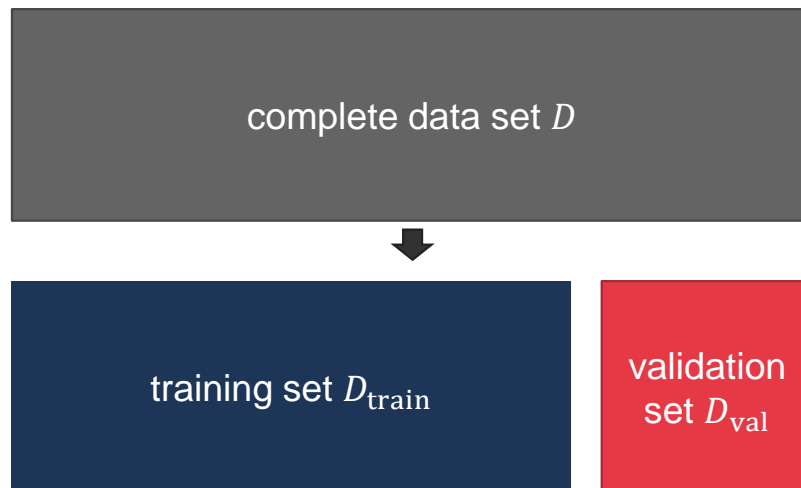


# Training and Validation Data Splitting



Cyber-Physical Systems  
in Mechanical Engineering TU Berlin

- **Aim:** find well-generalizing models with low bias and low variance
- **Generalization property:**
  - Generalization is the act of performing tasks of the same difficulty and nature
  - In data-driven modeling: predictive model performs well on previously unseen data



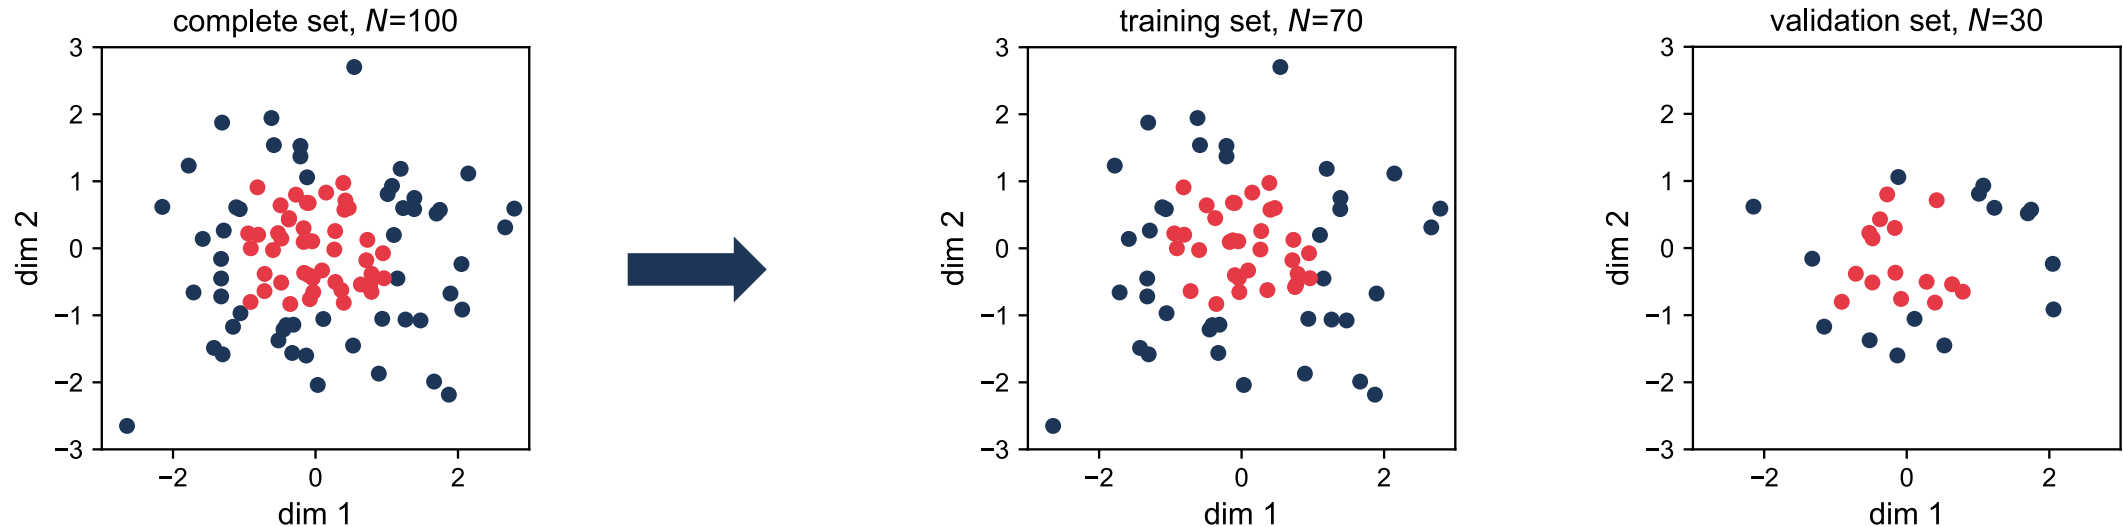
## Typical procedure in machine learning:

1. Split data into training and validation set (typically at ratio of 0.7 / 0.3 or 0.8 / 0.2)
2. Find best model parameters  $\theta$  on training set  $D_{\text{train}}$
3. Evaluate performance on hold-out validation set  $D_{\text{val}}$

# Example



- Train-validation split through random sampling (without replacement)



- Potential risk: different data distributions in training and validation set
  - Check data split distributions, particularly for small data sets
  - Consider using stratified splits, particularly for imbalanced data sets



# Measuring Classification Performance



Cyber-Physical Systems  
in Mechanical Engineering TU Berlin

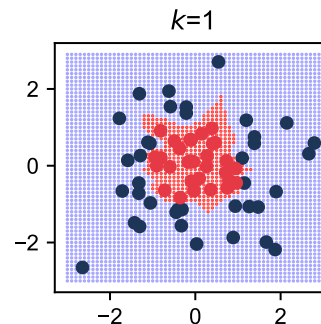
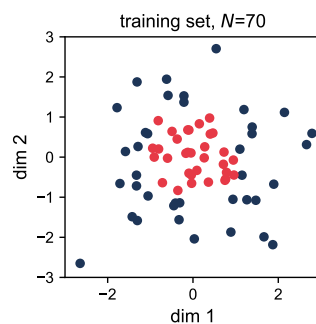
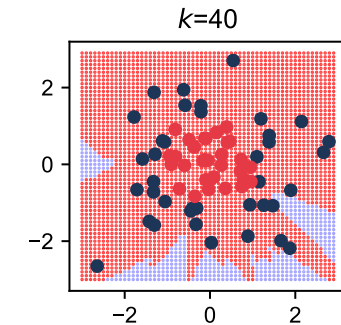
- To choose an appropriate  $k$  for kNN, we need some quality metric for how many of the data points got classified correctly
  - Straight-forward metric: **accuracy**  $\text{acc} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$  ratio of correctly predicted samples
  - **TP**: number of true positives  $y = 1, \hat{y} = 1$
  - **FP**: number of false positives  $y = 0, \hat{y} = 1$
  - **TN**: number of true negatives  $y = 0, \hat{y} = 0$
  - **FN**: number of false negatives  $y = 1, \hat{y} = 0$
- |              |           |
|--------------|-----------|
| ground truth | $y$       |
| prediction   | $\hat{y}$ |
- Positives/negatives: depends on the application and the perspective. Typically: positive class is the quantity you want to predict / detect.
  - **Example**: Structural health monitoring, prediction of cracks in concrete structures.
    - Positive class: crack exists
    - Negative class: no crack

Which are the most  
critical false predictions?

# Over- and Underfitting

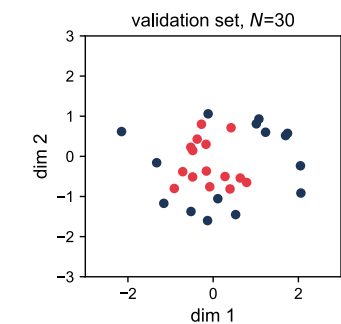


- **Underfitting:** prediction quality is low in training and validation data set
  - Model capability is not matching the prediction task complexity
- **Overfitting:** high prediction quality on training set, weak quality on validation set
  - Model is too complex for the given data set
  - Data set is too small / has not enough variance / is strongly imbalanced
  - Model has picked up all peculiarities of the training set, but no meaningful underlying patterns that are required for new unseen data



Performance on  
training set:  $\text{acc} = 1.0$

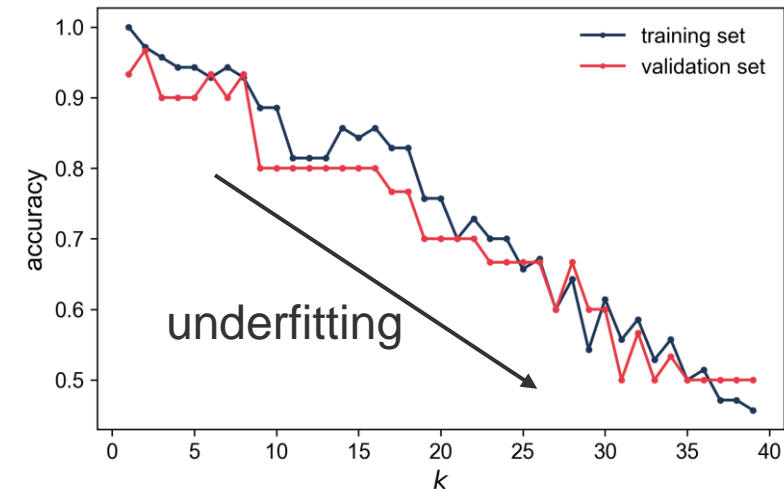
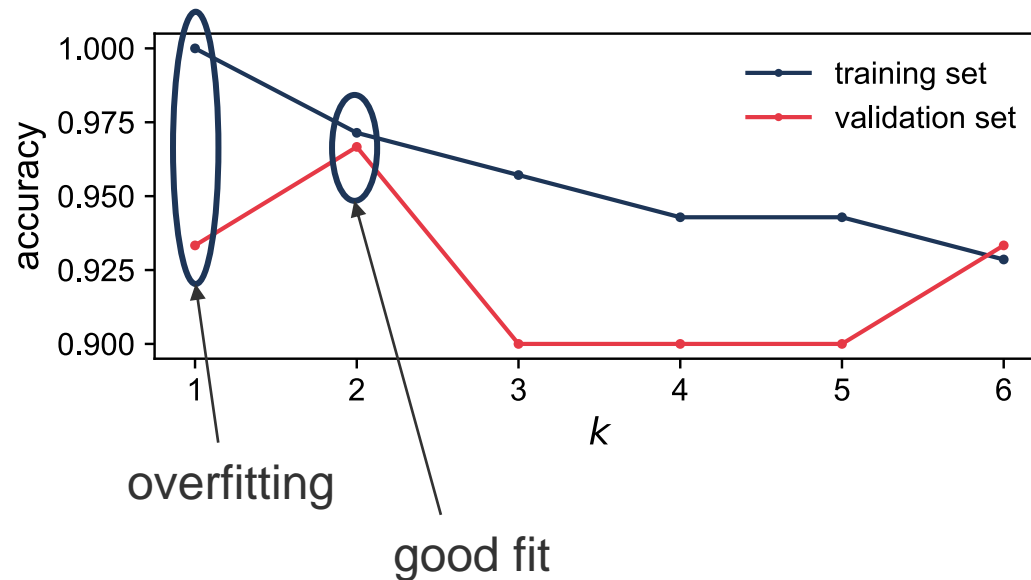
Performance on  
test set:  $\text{acc} = 0.933$



# kNN: choosing $k$



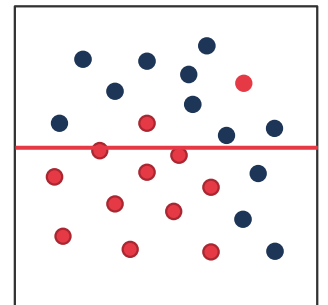
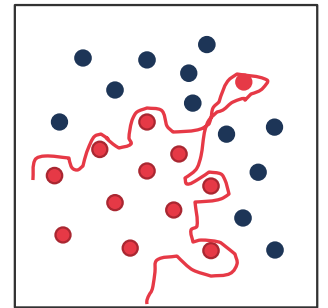
- Edge cases:
  - $k = 1$  accuracy 1.0 on training data set
  - $k = N$  accuracy  $\hat{N}/N$  on training data set,  $\hat{N}$  number of majority class members
- Different choices of  $k$ :  $k \in \{1, \dots, 40\}$



# Bias Variance Tradeoff



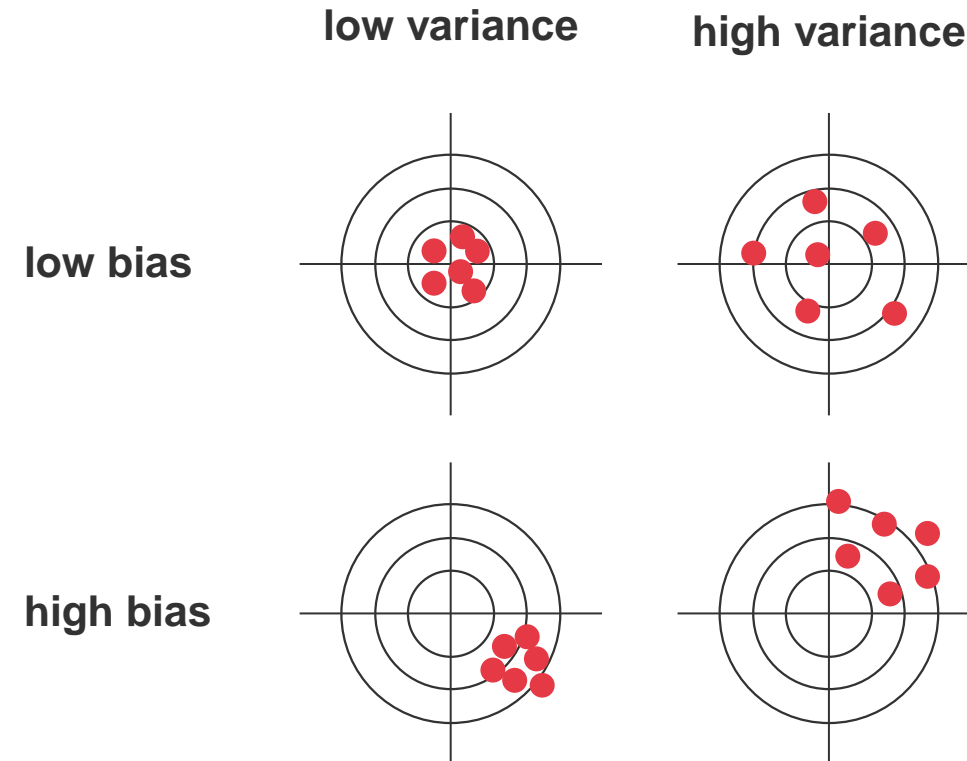
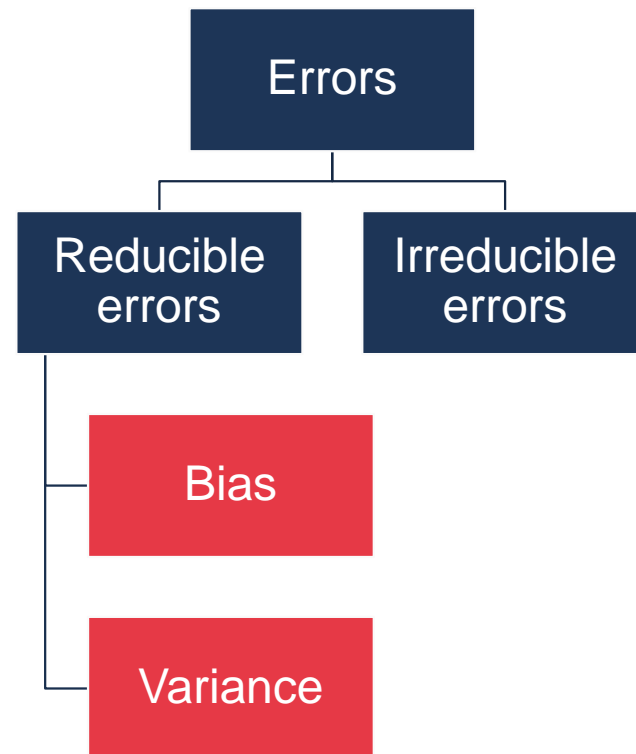
- Setting: train different versions of a given model on different training data sets (obtained from the same original set, e.g. through bootstrapping). Every version of the model will make different predictions on a hold-out validation set
- **Variance:** a measure how far each model's predictions are from the average across all models
  - Large variance: model versions perform vastly different on different data sets
  - Caused by high flexibility of the model, potentially over-adjusting to the specific training data set
  - Large variance → **overfitting**
- **Bias: Model unable to capture the true relationship between inputs and targets**
  - A measure how far the average model prediction is from the ground truth
  - Large bias: on average, all model versions deviate strongly from the ground truth
  - Caused by a model that is unable to capture the dominant patterns in the data
  - Large bias: → **underfitting**



# Bias Variance Tradeoff



- Bias and variance are **reducible errors** in data-driven modeling

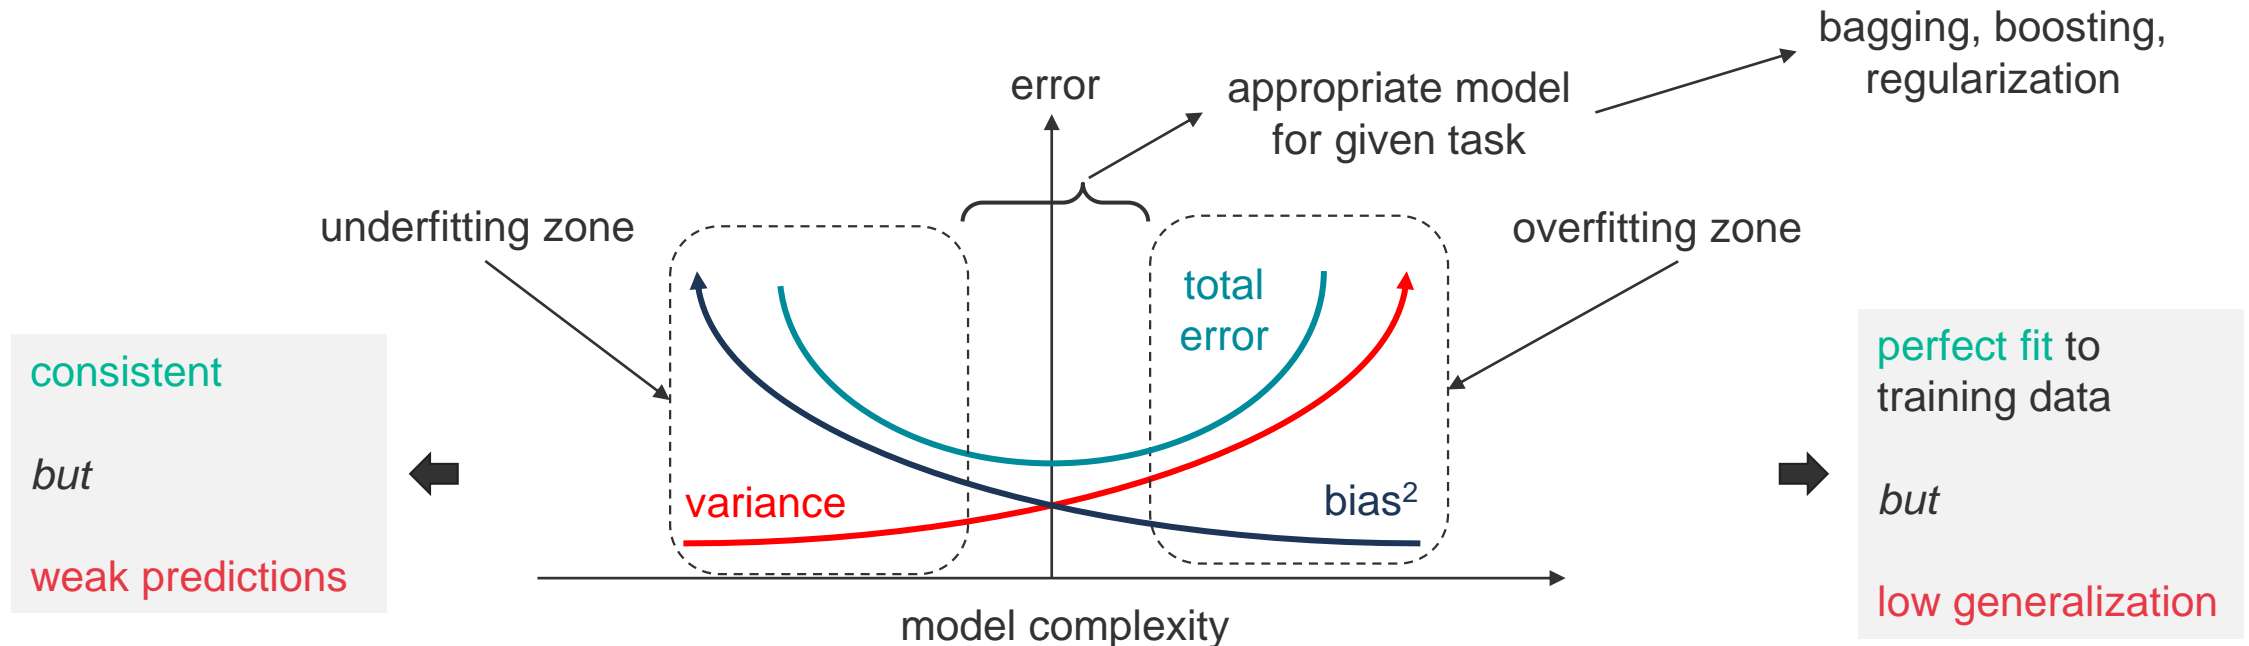


# Bias Variance Trade-Off



Cyber-Physical Systems  
in Mechanical Engineering TU Berlin

- Reducible errors, **but** there is a **tradeoff** between variance and bias



- A biased model can perform better on the validation data than a model that was overfitted to the training data, hence finding an optimal model can be a tradeoff

# Bias and Variance: How to Proceed



- Typical approach to machine learning modeling tasks
- 1. **Start with a very simple model** (*baseline* model) with high bias
- 2. **Increase model complexity until acceptable bias obtained**
- 3. **Validation set evaluation.** In case of high variance:
  - a. Try to get more data or better data (e.g. data augmentation)
  - b. Make changes to the model architecture (regularization)
  - c. Try different model architecture and continue with (2)

Reduce bias



Reduce variance

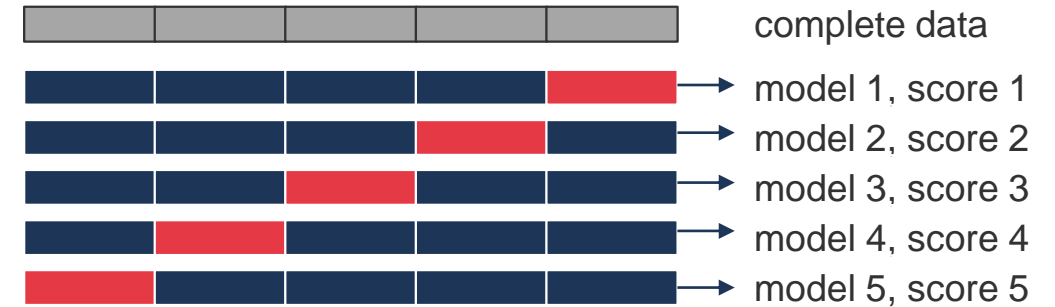
→ **Diagnose the problem first, then take actions!**

# k-fold Cross Validation

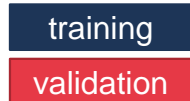


- **Procedure:**

1. Split data set into  $k$  subsets
2. Use  $(k - 1)$  subsets as **training set**, and remaining subset as validation set
3. Repeat by iterating over all subsets



- Report of evaluation metric:       $\text{score} = \text{mean} \pm \text{std. dev. (validation set score)}$



- **When to use?**

- Comparatively small data set with strong data spread
- Simple train-test split yields non-repeatable results (within some tolerance)
- Measure bias and variance statistics of your model
- Compare different model architectures against each other



# k-fold Cross Validation



## `sklearn.model_selection.KFold`

```
class sklearn.model_selection.KFold(n_splits=5, *, shuffle=False, random_state=None)
```

[\[source\]](#)

K-Folds cross-validator

Provides train/test indices to split data in train/test sets. Split dataset into k consecutive folds (without shuffling by default).

Each fold is then used once as a validation while the k - 1 remaining folds form the training set.

Read more in the [User Guide](#).

### Parameters:

**`n_splits` : int, default=5**

Number of folds. Must be at least 2.

*Changed in version 0.22:* `n_splits` default value changed from 3 to 5.

**`shuffle` : bool, default=False**

Whether to shuffle the data before splitting into batches. Note that the samples within each split will not be shuffled.

**`random_state` : int, RandomState instance or None, default=None**

When `shuffle` is True, `random_state` affects the ordering of the indices, which controls the randomness of each fold. Otherwise, this parameter has no effect. Pass an int for reproducible output across multiple function calls. See [Glossary](#).

# Cross Validated Choice of $k$



- Using 5-fold cross validation for selecting an optimal value of  $k$ . Validation scores:

$k$	mean accuracy	std. dev accuracy	bias / variance
1	0.94	0.058	low bias, high variance
<b>2</b>	<b>0.95</b>	<b>0.032</b>	<b>low bias, low variance</b>
3	0.93	0.051	low bias, high variance
4	0.94	0.049	low bias, high variance
5	0.90	0.550	high bias, high variance
10	0.91	0.037	high bias, low variance
20	0.78	0.040	high bias, high variance

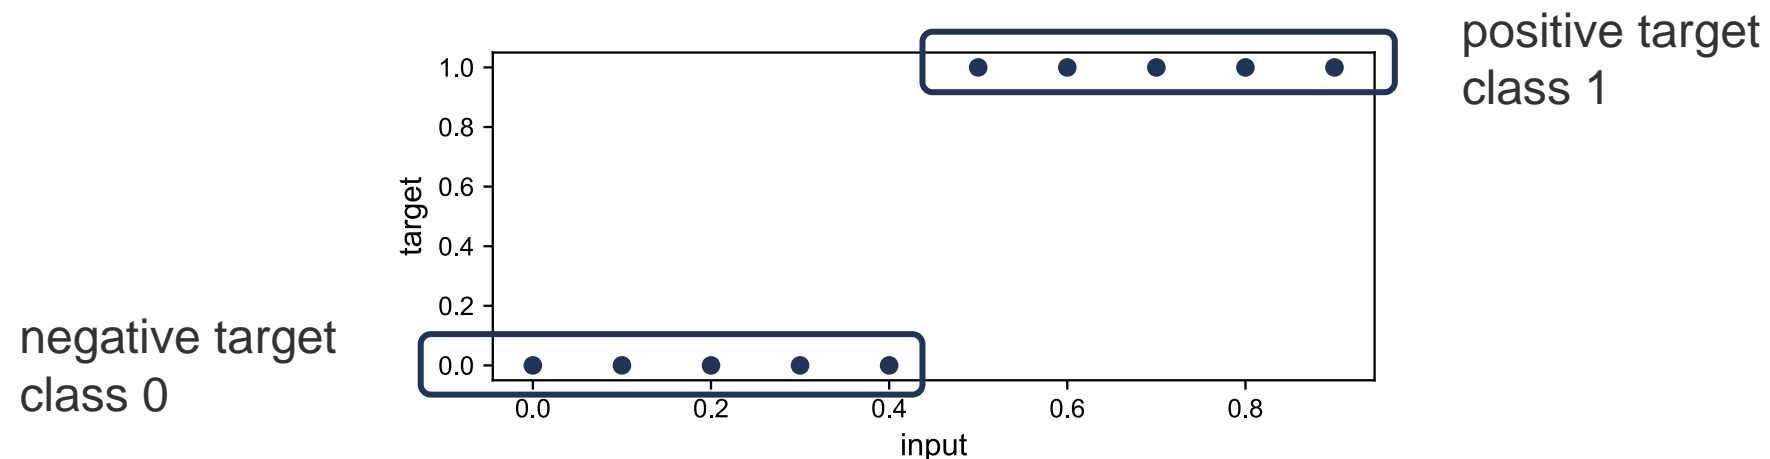


# Logistic Regression

# Logistic Regression



- Simple **classification method**
- Why 'regression'? We are in fact predicting probabilities  $[0, 1]$  and not categories
- Based on maximum likelihood
- Different names: *logit regression*, *maximum-entropy classification* or *log-linear classifier*

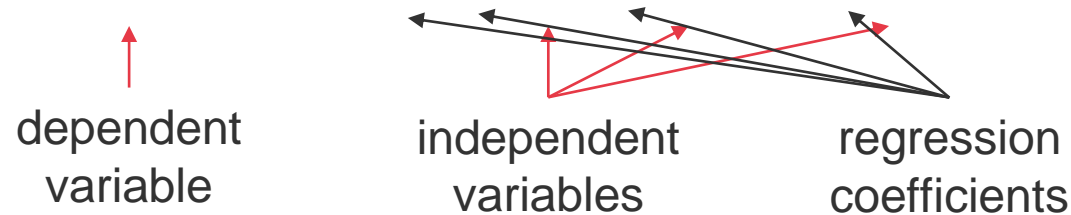


# Logistic Regression

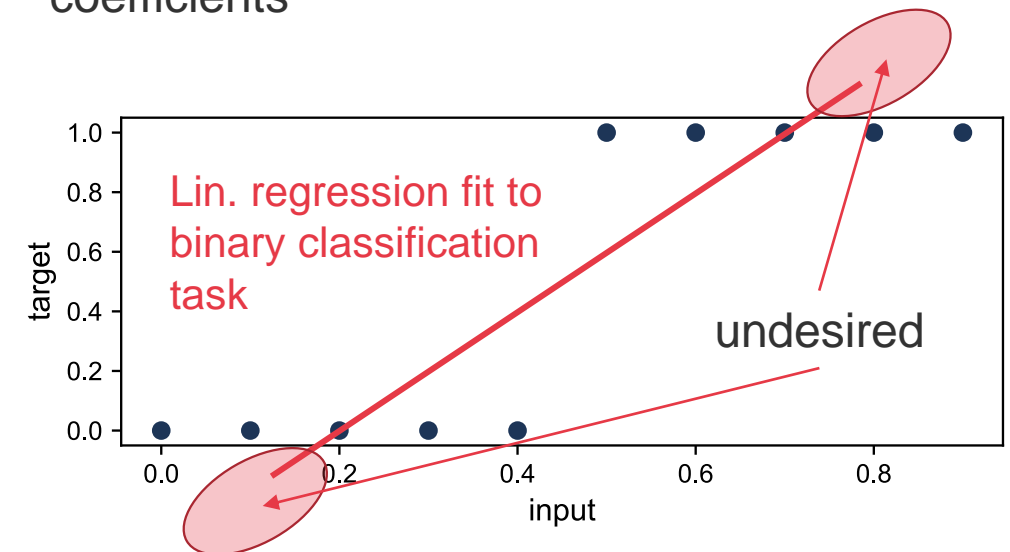


- Logistic regression: estimate the probability of observing a positive label given the current inputs

- Linear regression:**  $\hat{y}_i = \mathbf{x}_i^T \boldsymbol{\theta} = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots \theta_n x_n$



- Logistic regression:**
  - Dependent variable on nominal scale
  - Linear regression would estimate the value itself, not the probability of observing the positive class
  - Aim: prediction of  $P(y = 1 | \mathbf{x})$ , i.e. the probability of the positive (1) target



# Logistic Regression



- Logistic regression: estimate the probability of observing a positive label given the current inputs  $\mathbf{x}$

- Linear regression:**  $\hat{y}_i = \mathbf{x}_i^\top \boldsymbol{\theta} = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$

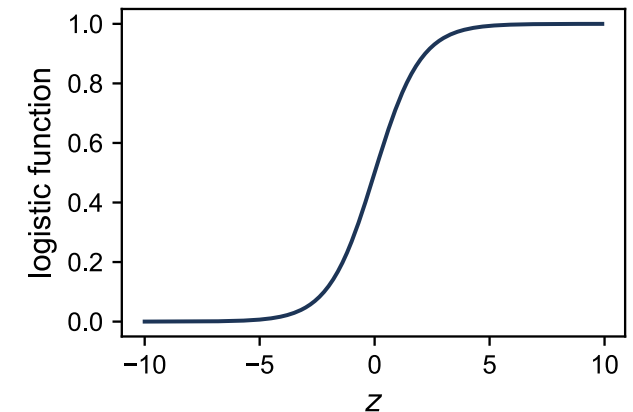
- Logistic regression:**
  - Restrict output range to  $[0, 1]$  using sigmoid

- Idea: take the logistic of the linear regression:

$$\hat{y} = \frac{1}{1 + e^{-(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n)}}$$

Logistic function (sigmoid function)

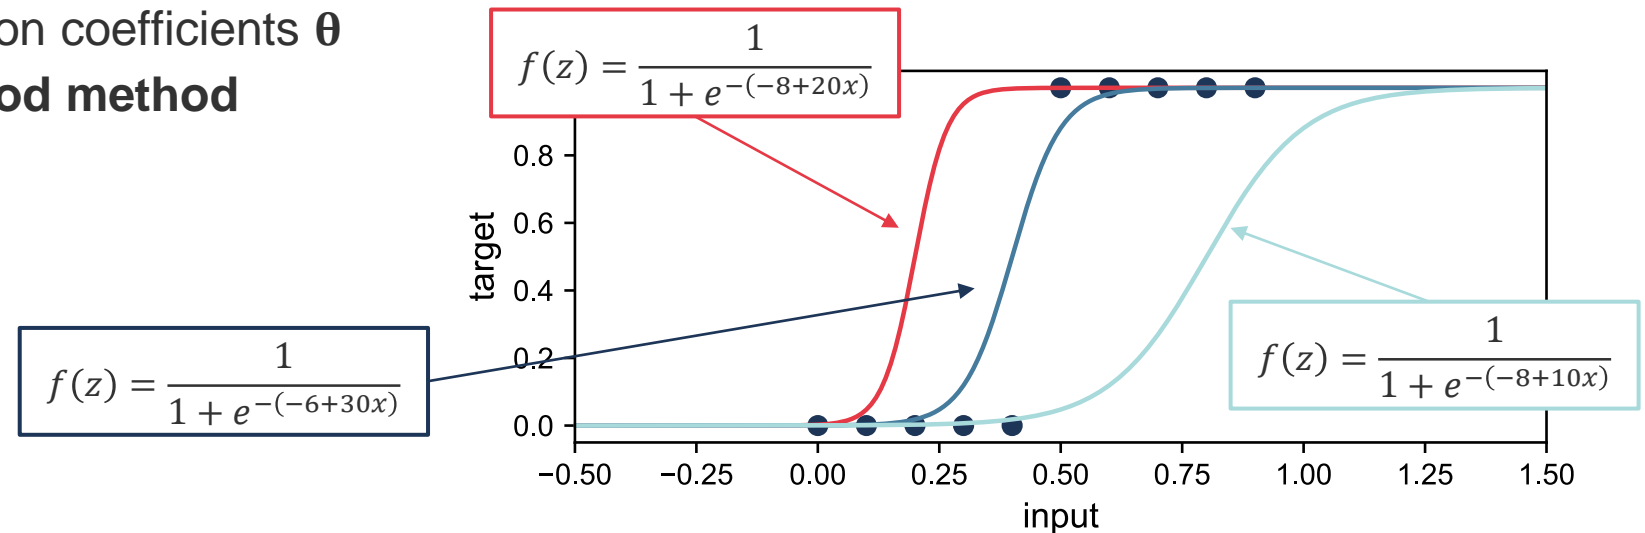
$$f(z) = \frac{1}{1 + e^{-z}} \in [0, 1] \quad \forall z$$



# Logistic Regression



- **Logistic regression:** 
$$\hat{y} = \frac{1}{1 + e^{-(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n)}}$$
- Probability of observing the positive class: 
$$P(y = 1 | x_1, x_2, \dots, x_n) = \frac{1}{1 + e^{-(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n)}}$$
- Probability of observing the negative class: 
$$P(y = 0 | x_1, x_2, \dots, x_n) = 1 - \frac{1}{1 + e^{-(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n)}}$$
- Computation of regression coefficients  $\theta$ 
  - **Maximum likelihood method**



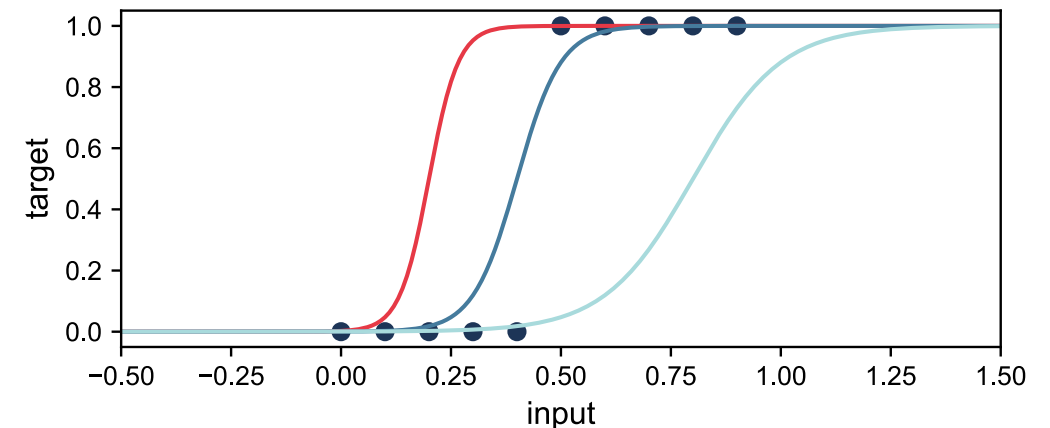
# Maximum Likelihood



- Likelihood function  $L(\theta)$  denoting the probability of observing the positive class for a given input  $x$
- **Maximum likelihood estimator:** maximize log likelihood  $LL(\theta) = \log(L(\theta))$

→ Minimize  $LL(\theta) = \sum_{i=1}^N (-y_i \log(\hat{p}(x_i)) - (1 - y_i) \log(1 - \hat{p}(x_i)))$

- **Methods for finding optimal  $\theta^*$ :** nonlinear optimizers
  - Stochastic gradient descend (SGD)
  - Newton-Cholesky
  - Newton-Conjugate Gradient
  - ...







# Exercise 08

May 31, 2023

# Exercise 08



- Implement k-nearest neighbors from scratch
  - Object-oriented implementation
  - .fit(X, y), .predict(X) methods
- Implement classification scores
  - Accuracy
  - Recall
  - Precision
- Evaluate tendency to under- and overfitting for
  - kNN,
  - decisionTrees
  - logistic regression



# Questions?